

ZHAO JIN

+1 404-668-1916 | lunamos.thu@gmail.com | zehaojin.com

EDUCATION

Georgia Institute of Technology, Atlanta, Georgia, US

Jan 2026 – May 2027 (expected)

Master of Science in Computational Science and Engineering (GPA: 4.0/4.0)

Tsinghua University, Beijing, China

Sep 2021 – Jul 2025

Bachelor of Science in Mechanics and Aerospace Engineering (GPA: 3.8/4.0)

PROFESSIONAL EXPERIENCE

StepFun

May 2026 – Present

Foundation Model Research Intern | Shanghai, China

- Ran SFT, RL, and OPD experiments for **coding agents**, achieving team-best results on **SWE-Pro** and **SWE-Atlas** with an internal model. The resulting training recipes are being incorporated into the next model iteration.
- Built an LLM-orchestrated pipeline to generate, refine, and evaluate **coding-agent environments**, with automated checks for task and rubric quality and the completeness of workspaces and attachments.
- Defined and validated 10+ trace-level indicators for **agent failure modes**, focusing on long-horizon context inefficiency; benchmarked them across frontier coding models for behavioral analysis and failure attribution.

Meituan LongCat LLM Team

Sep 2025 – Dec 2025

Foundation Model Research Intern | Shanghai, China

- Post-trained Qwen3-VL-32B for physical-world reasoning, improving its ERQA score by 4 points.
- Built the **VLA training and evaluation** stack for pi0, pi0.5, and Isaac GR00T on LIBERO, and reproduced end-to-end pretraining and fine-tuning pipelines for the pi model family.
- Developed a memory-augmented **VLM gaming agent** for Atari in the Arcade Learning Environment, enabling GPT-4o to improve through iterative experience summarization and achieve near-human scores on 3 games.
- Co-first-authored **3 papers** on **multi-turn instruction following** and **real-world business evaluation**, including SOP-Maze (ACL Findings 2026), with additional submissions to EMNLP and AACL.

RESEARCH HIGHLIGHTS

Beyond Steering Vector: Flow-based Activation Steering for Inference-Time Intervention

Mar 2026 – May 2026

Georgia Institute of Technology | Atlanta, Georgia

- Proposed FLAS (Flow-based Activation Steering), a concept-conditioned velocity field for **LLM inference-time intervention** that models multi-step, token-varying activation trajectories.
- Achieved **state-of-the-art** zero-shot steering on AxBench, becoming the first steering method to consistently outperform in-context prompting; outperformed HyperSteer by 67% using less than 4% of its parameters.
- Project Page: <https://flas-ai.github.io>; ICML 2026 MechInterp workshop poster; under review at NeurIPS 2026.

Stochastic Attention: Randomized Routing for Expressive Linear-Time Attention

Dec 2025 – Apr 2026

Tsinghua University | Beijing, China

- Proposed Stochastic Attention (SA), a parameter-free enhancement for sliding-window attention by introducing random token permutations. Achieved up to **28× speedup over full attention** at 32K sequence length.
- Pretrained 360M-parameter language models** on SlimPajama with a gated SA+SWA dual-path architecture, outperforming prior state-of-the-art methods in average zero-shot accuracy across seven benchmarks.
- Paper link: <https://arxiv.org/abs/2604.00754>, under review.

SOP-Maze: Evaluating LLMs on Complicated Business Standard Operating Procedures

Sep 2025 – Dec 2025

Meituan LongCat LLM Team | Shanghai, China

- Proposed SOP-Maze, the first benchmark for **evaluating LLMs on real-world business SOPs**, comprising 397 instances and 3,422 subtasks across 23 complex scenarios.
- Developed a deterministic evaluation framework based on JSON Schema and reference-index matching, enabling reliable and efficient scoring of free-form model responses.
- Paper link: <https://arxiv.org/abs/2510.08942>, ACL 2026 Findings.

SELECTED HONORS AND AWARDS

- Ying-Hua Fellowship**, Tsinghua University (for 15 undergrads) Dec 2022
- Spark Fellowship**, Tsinghua University (for 40 undergrads) May 2023
- Overall Excellence Scholarship**, Tsinghua University (for top 5% undergrads) Nov 2022 & 2024
- Finalist**, Interdisciplinary Contest in Modeling (ICM) (for top 2% teams) May 2023

SKILLS

Language: Fluent in English (TOEFL 109) and Mandarin.

Technical: PyTorch, Hugging Face Transformers, Megatron, vLLM, Python, C/C++.

Research Focus: Foundation Model Post-Training, AI Agents, Evaluation, and Mechanistic Interpretability.